

# Algoritmi di Crittografia

## Corso di Laurea Magistrale in Informatica

A.A. 2018/2019

1

## Randomness

- Basic probability
- Random (and Pseudo-Random) Number Generators

# Algoritmi di Crittografia

1

## Randomness

- Basic probability
- Random (and Pseudo-Random) Number Generators

# Randomness and cryptography

- Modern crypto applications would be simply not possible without the availability of some source of *randomness*
- For instance, randomness is required to generate:
  - asymmetric keys in RSA-based protocols;
  - symmetric keys in SSH sessions;
  - stored salt data to be hashed with passwords;
  - initialization vectors to a block cipher's mode of operation.
- As we shall see, the ideal source of randomness for these (and other crypto) applications produces sequences of bits that are: (1) *independent*, (2) *uniformly distributed*, (3) always ready to hand
- Not as easy as it might sound...
- But let's start with some basic probability

# Probability (informal)

- The probability of an event is a measure of “how likely is the occurrence of that event”
- Such measures are normalized values in the interval  $[0, 1] \subseteq \mathbb{R}$
- Over discrete set of events, 0 means “impossible” while 1 means “certain”
- Over the reals (e.g., when dealing with time) things are little more complicated
- In computer science applications (and hence in Cryptography) the sets of events are always discrete
- But what exactly is the probability of an event? How can we assign such numbers to the events that might occur?
- Not always easy an answer

# Probability (informal)

- Start by considering an experiment involving randomness (e.g., tossing a coin, rolling a dice, ...)
- Each possible outcome (e.g., the face which the coin ends up showing) is an *elementary event*
- Each elementary event is given (here is where we can make mistakes ...) a probability, i.e., a real number in  $[0, 1]$
- The constraint is that the sum (finite or infinite) of the probabilities to all the events must be 1
- An *event* is then any subset of the set of all the elementary events
- Example, “the dice ends up showing a face with an even number of dots” is an event composed of 3 elementary events
- The probability of an event is the sum of the probabilities of the corresponding elementary events

# Probability (decently formal)

- To reason about probability we need the following ingredients
- We are given a set  $\Omega$ , finite or countably infinite, called the *sample space*
- We consider a function (a *measure*) over  $2^\Omega$  with values in  $[0, 1]$ ; call them  $p$  (for probability...)
- $p$  is defined for any  $A \subseteq \Omega$  (i.e.,  $A \in 2^\Omega$ )
- In particular,  $p$  is defined for any  $\omega \subseteq \Omega$  such that  $|\omega| = 1$  ( $\omega$  is an elementary event)
- For any  $A \subseteq \Omega$ , it holds that

$$p(A) = \sum_{\omega \in A} p(\omega)$$

- If  $A$  is empty,  $p(A) = 0$  while, if  $A = \Omega$ ,  $p(A) = 1$

# Example

- We consider the experiment of tossing a coin and denote with *head* and *tail* the corresponding elementary events
- $\Omega = \{head, tail\}$
- Assuming the two outcomes are equally likely, and that it is impossible for the coin to land on the edge, we can set:

$$p(head) = p(tail) = \frac{1}{2}$$



# Probability distributions

- The “form” of the probability function is known as the (probability) *distribution*
- When all the elementary events have the same probability (which implies that the sample space is finite) we speak of *uniform* distribution
- The distribution of the coin tossing example of the previous slide is uniform
- Uniform distributions are the most important in cryptographic applications, but many different (and important) distributions exist
- We will see one in the next two slides

# Poisson distribution

- Suppose you (as a web server admin) are interested in the following question “How many http requests will arrive in the next  $t$  seconds”?
- Clearly, even if we may sensibly assume that the number of requests will be a finite quantity, we cannot place an à priori bound
- Hence, we want to assign a probability to all the events: no request (in  $t$  seconds), just one request, exactly two requests, and so on
- To do this, we need a model, which can be more or less accurate
- We assume to know the *rate* with which http requests arrive, a quantity that we denote by  $\lambda$

# Poisson distribution

- Let us define:

$$p(k \text{ requests in } t \text{ secs}) = \frac{(\lambda t)^k}{k!} e^{-\lambda t}$$

- Since  $p(k \text{ requests in } t \text{ secs}) \geq 0$  and

$$\sum_{k=0}^{\infty} p(k \text{ requests in } t \text{ secs}) = \sum_{k=0}^{\infty} \frac{(\lambda t)^k}{k!} e^{-\lambda t} = 1$$

it follows that  $p$  is a valid probability function

- Under reasonable assumptions, it is also a good choice, meaning that it produces accurate estimates of the process under analysis

# Random variables

- Random variable represent a very useful way of dealing with sample spaces
- Consider the experiment of rolling a couple of dice
- What are the elementary events here?
- Clearly, a dice rolls and (when it ends rolling) the important fact is the number of dots in its up-facing side
- So, to describe one possible outcome, we should rigorously (but quite tediously) write something like:

*The first dice ended up rolling with the up-facing side showing 3 dots and the other dice ended up rolling with the up-facing side showing 2 dots*

- Why not to say simply: one (dice) is 3 and the is 2?
- Because 3 and 2 are numbers, not events
- Here is where random variables come up

# Random variables

- A random variable maps events to a numeric space  $S$
- In  $S$  it is easier to make quantitative reasonings
- Random variables also make it possible to describe events in an easier way
- Let  $X$  be the random variable that maps the 6 possible outcomes (elementary events) of the experiment of rolling a dice
- $X = 2$  is then a (very convenient) way to denote the elementary event: “The dice ended up rolling with the up-facing side showing 2 dots”
- Similarly,  $x \leq 3$  denotes the event: “The dice ended up rolling with the up-facing side showing 1, 2, or 3 dots”

# Random variables

- If  $X$  is a random variable and  $x \in S$ , then the probability that  $X$  has the value  $s$ , denoted  $p[X = s]$ , is

$$p[X = s] = \sum_{\omega \in \Omega | X(\omega) = s} p(\omega)$$

i.e., is the sum of the probabilities of all the elementary events that maps to  $s$

- Let  $X$  be a random variable that maps the outcomes of rolling two dice to  $S = \{2, 3, \dots, 12\}$  in the obvious way, namely:

$$X(\text{1 dot}, \text{1 dot}) = 2$$

$$X(\text{1 dot}, \text{2 dots}) = 3$$

...

$$X(\text{6 dots}, \text{6 dots}) = 12$$

# Random variables

- But then for instance:

$$p[X = 4] = p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare)$$

- We can also define the value of  $X$  using a richer mathematical armory:

$$p[X \leq 3] = p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare)$$

or

$$\begin{aligned} p[X > 7 \text{ and } X \leq 10] = & p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + \\ & p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + \\ & p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + \\ & p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) + p(\blacksquare, \blacksquare) \end{aligned}$$

# Independence

- This is the last notion we recall explicitly, given its importance in cryptography
- Consider two events,  $E_1$  and  $E_2$ , not necessarily defined over the same sample space
- We say that  $E_1$  and  $E_2$  are independent events if the occurrence of  $E_1$  does not affect the probability of  $E_2$  occurring, and vice versa
- If we toss a coin and roll a dice, the probability that the coin lands heads is  $\frac{1}{2}$  (assuming the coin is not biased), independently of the outcome of the dice experiment



# Independent random variables

- Consider rolling two six-face dice and assume that the dice are not biased, so that we can fairly assume that all the possible outcomes have probability  $\frac{1}{6}$ .
- Let  $X$  and  $Y$  be random variables that refers to the first and second dice experiment, respectively
- Then,  $X$  and  $Y$  being independent amounts to the following property:

$$p(X = x \text{ and } Y = y) = p(X = x) \cdot p(Y = y) = \frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$$

where clearly  $x, y \in \{1, 2, 3, 4, 5, 6\}$

- If  $X$  and  $Y$  are independent we may also write, e.g.:

$$p(X \leq 4 \text{ and } Y > 2) = p(X \leq 4) \cdot p(Y > 2) = \frac{4}{6} \cdot \frac{3}{6} = \frac{1}{3}$$

# Algoritmi di Crittografia

1

## Randomness

- Basic probability
- Random (and Pseudo-Random) Number Generators

## A first simple example

- For cryptographic applications, it is crucial to have sources of randomness, i.e., logical or physical devices that can produce (measurable) random events in some space
- What ultimately matters is to map the available randomness to sequences of *independently generated bits* (or, simply, *independent bits*)
- Suppose we need a randomly chosen binary sequence of length  $n$  (say, to generate a session key as part of some symmetric protocol)
- All the possible  $2^n$  sequences must have the same chance to be chosen
- Assume also, for simplicity, that  $n = 4$

## A first simple example

- Suppose to this end we use a Random Bit Generator, i.e., a piece of software that, each time is called, returns 1 with probability  $p$  and 0 with probability  $1 - p$  independently of *past history*
- In other words, the bits are independently generated
- Our question is: what is the probability of generating (say) 1011 as key?
- Well, this is easy: since the bits are independent, the probability is  $p(1 - p)pp = p^3(1 - p)$
- In general, a key with  $k$  ones and  $4 - k$  zeroes will show up with probability:

$$\binom{4}{k} p^k (1 - p)^{4-k}$$

## A first simple example

- Now we ask a different question, with a clear cryptographic flavor: how hard would be for  $\mathcal{E}$  to guess our session key?
- Suppose  $p = 1$  and (clearly)  $1 - p = 0$ . In this case only one key is possible, namely 1111 which requires no guess at all...
- Suppose now that  $p = 0.9$ . All the 16 4-bit long sequences have their chance now. But the sequence 1111 is still the best guess for  $\mathcal{E}$ , who will be right approximately 66% of the times
- So, what is the value of  $p$  that cause  $\mathcal{E}$  having a hard time to guess correctly?

# Surprise and ... information

- Let's analyze things from another perspective
- Assume  $\mathcal{E}$  knows the value of  $p$  and consider again the state of affairs where  $p = 1$ .
- In this case releasing to  $\mathcal{E}$  the “information” that the key is 1111 actually amounts to giving her no-info at all
- Suppose now that  $p = 0.9$ . This time releasing the same info as above would indeed carry some information, however one that definitely would not knock Eve out
- It then seems that the question raised in the previous slide can be rephrased as follows: what is the value of  $p$  that implies maximum gain of information from the outcome?

# Entropy

- Entropy is a well-known concept in thermodynamics, which is however not relevant to us
- In *Information theory*, the entropy of a random event is a measure of the average information that one acquires when the outcome is revealed
- ... and this is exactly what we need
- If the event has  $k$  possible outcomes, and  $p_1, p_2, \dots, p_k$  denote their probabilities, then the entropy is defined as

$$\sum_{i=1}^k p_i \log_2 \frac{1}{p_i} = - \sum_{i=1}^k p_i \log_2 p_i$$

- The values we are looking for are then those  $p_i$  that maximize the above quantity, under the constraints:  $p_i \geq 0$  and  $\sum_{i=1}^k p_i = 1$

# Entropy

- If  $k = 2$ , hence  $p_1 = p$  and  $p_2 = 1 - p$ , it is an easy exercise to prove that entropy is maximum for  $p = \frac{1}{2}$
- The result carries over the general case, namely  $p_i = \frac{1}{k}$ , for  $i = 1, \dots, k$ , are the values that implies maximum entropy
- For our toy example, it  $p = \frac{1}{2}$ , then all the keys have the same probability of  $\frac{1}{16}$  and the entropy is

$$16 \left( -\frac{1}{16} \log_2 \frac{1}{16} \right) = -\log_2 \frac{1}{16} = \log_2 16 = 4$$

- In other words, in case of  $p = \frac{1}{2}$  we learn four bits of information and this is supposed to be maximum information
- Indeed, it makes sense: how would it be possible to get more than four bits of information when the key is exactly four bit long?



# Entropy

- Any other (i.e., non-uniform) distribution gives lower entropy
- The same happens if the bits are not independent
- For instance, suppose the probability of the next bit being 1 is  $1/4$ ; then the entropy  $E$  associated to the next two bits is:

$$\begin{aligned} E &= -\text{prob}(00) \log \frac{1}{\text{prob}(00)} - \text{prob}(01) \log \frac{1}{\text{prob}(01)} - \\ &\quad \text{prob}(10) \log \frac{1}{\text{prob}(10)} - \text{prob}(11) \log \frac{1}{\text{prob}(11)} \\ &= \left(\frac{3}{4}\right)^2 \log \left(\frac{4}{3}\right)^2 + 2 \frac{3}{4^2} \log \frac{4^2}{3} + \left(\frac{1}{4}\right)^2 \log 4^2 \\ &\approx 1.62 \end{aligned}$$

# Random Number Generators (RNG)


- The ideal source of randomness would then be a Random Bit Generator which always returns 0 or 1 with uniform probability (i.e., independently of past bits)
- In general, in the crypto literature we speak of *Random Number Generators*, rather than of random bit generators, irrespectively of the actual device being oriented to bits or numbers
- Aside for terminology, to generate truly random bits is all but easy in practice
- For non-crypto applications that embody a probabilistic model of some sort, *pseudo-randomness* is almost always adequate, and pseudo-randomness is much easier to obtain
- For crypto application pseudo-randomness is almost unavoidably the origin of disasters

# The problem in a nutshell

- Having random bits ready at hand (and at will) when they are needed is problematic
- True randomness can be extracted by the physical environment (including quantum mechanical phenomena)
- There might be availability problems (we need physical devices attached to the computer)
- Also, it might not be easy to transform the environmental randomness in a perfectly uniform and independent sequence of bits
- Gathering randomness from the computer itself (e.g., from disk activity, keyboard or mouse inputs) is clearly easier but might not be completely secure
- In all cases, the rate at which the random bits can be produced is not sufficient in many application settings (e.g., think of Amazon web servers...)

# A general solution

- Combine true RNGs with good Pseudo-Random Number Generators (*PRNGs*)
- Essentially, the RNG is used to “seed” a cryptographic PRNG (and later to collect other randomness, if and when available)
- A cryptographic hash function or a cipher is then used to return pseudo-random data (the output) and to update the internal state



Simple PRNG operations


# Outline

- This “architecture” prevents an attacker from computing the internal state/seed from the output
- This is usually sufficient to prove *computational security* (rather than *unconditional security*)
- We shall describe three different (allegedly...) *Cryptographically Secure PRNGs (CSPRNGs)*
  - 1 The *ANSI X9.17* standard, a PRNG which was widely adopted especially in banking applications
  - 2 The *Fortuna* PRNG, used in Windows® systems
  - 3 The `/dev/random` and `/dev/urandom` generators under Linux and other \*nix systems

# The ANSI X9.17 generator

- Requirements:
  - A secret key  $k$  generated at initialization time
  - A *Triple-DES* symmetric block-cipher
  - A short initial random seed  $S$  (64 bits), which is (part of) the internal state
- When pseudorandom bits are required:
  - 1 Compute a timestamp  $ts$  of current date and time  $T$  at maximum possible resolution (e.g., microseconds)
  - 2 Using the encryption function of 3DES with key  $k$ :
    - compute  $T = E_k(ts)$
    - compute the output value  $O = E_k(T \oplus S)$
    - update the internal state:  $S = E_k(T \oplus O)$
  - 3 Return  $O$

# ANSI X9.17 block diagram



## A Python code for ANSI X9.17

```
from datetime import datetime
from Crypto.Cipher import DES3
from Crypto import Random
from Crypto.Util.strxor import strxor

class ANSIX917:
    def __init__(self, keylen):
        assert(keylen == 24 or keylen == 16)
        key = Random.new().read(keylen)
        self.__state = Random.new().read(keylen)
        self.cipher = DES3.new(key, DES3.MODE_ECB)
```



## A Python code ANSI X9.17 (cont.d)

```
class ANSIX917:
```

```
...
```

```
def __iter__(self):
```

```
    return self
```

```
def __next__(self):
```

```
    fmt = '%y-%m-%d_%H:%M:%S.%f '
```

```
    ts = datetime.now().strftime(fmt)
```

```
    T = self.cipher.encrypt(ts)
```

```
    out = self.cipher.encrypt(strxor(T, self.__state))
```

```
    self.__state = self.cipher.encrypt(strxor(T, out))
```

```
    return out
```

# Usage

```
In [1]: from ANSIX917 import ANSIX917
```

```
In [2]: PRNG = ANSIX917(keylen=16)
```

```
In [3]: next(PRNG)
```

```
Out [3]: b'\x8c\x98\xbf~\xf3\xae\xabrB\x87\xd8\xe7\x
```

```
In [4]: next(PRNG)
```

```
Out [4]: b'I0!\xaa:!\x8b\xe2d\xf34\xa3a\x19\xc0*|\x1
```

# Security issues

- At the macroscopic level the goals are obvious: it should be impossible for an attacker to recover both past and future generated bits
- Being able to protect the past is known as *backtracking resistance*
- Success in protecting future bits is the property of *prediction resistance*
- For backtracking resistance the general idea is to use *one-way* transformations: either a cryptographic hash function, which is one-way, or a cipher, which is not invertible without the key
- For prediction resistance, a good strategy is to regularly update the internal state with fresh random bits, a solution that the ANSI X9.17 PRNG does not implement

# Attacks to PRNGs

- We have a *direct cryptanalytic attacks* when the attacker is able to distinguish the outputs from truly random bists
- If the attacker, at some time, has access to the internal state of the PRNG (for whatever reason), then s/he will try to exploit this knowledge for as long as possible with *state compromise extension attacks*
- Finally, cryptanalytic attacks can take advantage of the attacker being in some control of (possible) PRNG inputs; in this case we speak of *input-based attacks*
- We now present a simple case study with the hope to shed light upon some of the above points

## A known attack to ANSI X9.17 PRNG

- Technically, it's a state compromise extension attacks
- However, it also exploits the low entropy in the PRNG inputs (and hence it can also be considered an input-based attack)
- This attack is directed towards discovering the ANSI X9.17 state using the *meet-in-the-middle* attack technique
- It works under the assumption that the attacker has somehow learned the key  $K$  (due to a flaw in the implementation, a leakage of information, ...)
- The attacker needs also two consecutive outputs from the PRNG

# A known attack to ANSI X9.17 PRNG

- The attack works as follows.
- We know that

$$\begin{aligned}T_i &= \mathbf{E}(K, t) \\ O_i &= \mathbf{E}(K, T_i \oplus S_i) \\ S_{i+1} &= \mathbf{E}(K, T_i \oplus O_i)\end{aligned}$$

where  $t$  is current *timestamp* while  $O_i$  is the  $i$ -th output.

- It follows that  $S_{i+1}$  can be computed in two different ways, namely:

$$S_{i+1} = \mathbf{E}(K, T_i \oplus O_i) \tag{1}$$

and

$$S_{i+1} = \mathbf{D}(K, O_{i+1}) \oplus T_{i+1} \tag{2}$$

## A known attack to ANSI X9.17 PRNG

- The point is that we (the alleged attackers...) do not know at which time the outputs and the new states were computed, i.e., which timestamps  $t$  were used
- This is, however, a low entropy input to the PRNG
- If we can approximate  $t$ , say within a second or few seconds (which is highly reasonable), then we are left with about  $\log_2 10^6 \approx 20$  bits of entropy, since we assumed resolution of microseconds
- This is well affordable for any commodity PC so that we can mount the following search:
  - 1 compute (1) and (2) for different values of  $t$
  - 2 stop when (and if), for some  $t_1$  and  $t_2$ :

$$\mathbf{E}(K, t_1 \oplus O_{i+1}) = \mathbf{D}(K, O_{i+1}) \oplus t_2$$

## A Python code for the attack (1)

```
def timestamps(T):  
    def advance():  
        from datetime import timedelta  
        nonlocal T  
        T += timedelta(microseconds=1)  
        return T.strftime(fmt)  
    for i in range(1000000): yield advance()  
  
def PSCAttack(ANSIX917gen, key):  
    fmt = '%y-%m-%d_%H:%M:%S.%f'  
    cipher = DES3.new(key, DES3.MODE_ECB)  
    t = datetime.now()  
    o1, o2 = next(ANSIX917gen), next(ANSIX917gen)  
    L = [] # A list of candidate seeds  
    ts = timestamps(t) # A generator of timestamps
```




## A Python code for the attack (2)

```
from Crypto.Util.strxor import strxor
from bisect import bisect_left
for j in range(100000): # Should be sufficient
    t = next(ts)
    T = cipher.encrypt(t)
    s1 = cipher.encrypt(strxor(T,o1))
    L.insert(bisect_left(L,s1),s1)
    d = cipher.decrypt(o2)
    s2 = strxor(d,T)
    idx = bisect_left(L, s2)
    if idx!=len(L) and L[idx]==s2:
        print("Success ' ' )
        break
#####
```

# A more sophisticated design: The Fortuna CS-PRNG

- So named after the Roman goddess of chance
- Fortuna's internal state is a pair  $(K, C)$ , where  $K$  is a 256-bit key while  $C$  is a 128-bit counter
- The generator works as depicted below:



Output and key update in Fortuna.  $E$  is the AES (or an AES-like) encryption function.

# Fortuna's detailed operations

**Initialization** Assign initial values to the counter  $C$  and the key  $K$

**Reseeding** Update the internal state (both  $C$  and  $K$ )

**Block generation** Using the AES block cipher, generate a number of 128-bit (16 bytes) blocks sufficient to satisfy the user's request

**Random data generation** Generate the number of random data requested by the user

**Collection** Collect real random data from entropy sources, used to reseed the generator

# Initialization and reseeding

- Initialization is exceedingly simple

$K = \text{b}'0'$  # *i.e. bytes*

$C = 0$

**return** (K,C) # *State of generator*

- Reseeding mixes the current key with (fresh random) data  $S$
- Fortuna uses double SHA256 hashes (which simply means SHA256 twice)

$K = \text{SHAd256}(K || S)$  # *// here is string/bytes concat*

$C = C + 1$

# Block generation

- The input is the number  $k$  of required blocks
- State  $(C, K)$  is clearly available (e.g., global or instance variables/attributes)
- $E$  is the AES encryption function

```
def GenerateBlocks(k)
    out = b''
    for i in range(k):
        b = E(K,C)
        out = out || b
        C += 1
    return output
```

- Block generation is meant as an internal function

# Random data generation

- The input is the number  $n$  of random bytes required
- $n$  must not exceed  $2^{20}$  (i.e., at most 1MB of data can be generated per-request)
- The reason is that increasing the size would also increase the statistical divergence from true randomness

```
if n>2**20:  
    raise some error  
B = GenerateBlock(ceil(n/16))[:n]  
K = GenerateBlock(2) # Update the key  
return B
```

# Collection

- This is the most important (and difficult) aspect of a CSPRNG
- Collecting real randomness from the environment requires (if not dedicated hardware, at least) special enablement of various drivers
- From the perspective of the CSPRNG, randomness collection thus means *storing* and *making available* the entropy collected from external sources (e.g., the drivers)
- Typically, entropy is gathered from events like mouse movements and keystrokes as well as other various timing sources
- In Fortuna, each source distributes the entropy it collects (evenly) over 32 different pools
- In this way, unless the attacker has full control of all the sources, s/he cannot reconstruct the contents of each pool

# Reseeding reviewed

- Reseeding clearly uses data coming from the pools
- Reseeding is done when pool  $P_0$  has enough data
- Reseeding number  $k$  uses data from all the pools  $P_j$  such that  $2^j$  divides  $k$
- Thus, for instance, pool  $P_0$  is used at each reseeding while  $P_2$  is used at reseeding 4, 8, 12 and so on
- This results in a sort of frequency/amount of entropy tradeoff for the attacker who snatched the internal state
- Suppose s/he does not control source 0: then the system is reseeded too frequently with data s/he does not know
- Suppose now the attacker does not control source 1 (or 2); then the generator is reseeded less frequently with data s/he does not know, but in much larger quantity



# PyCrypto implementation of Fortuna

```
from Crypto.Random.Fortuna import FortunaGenerator
from Crypto.Util.py3compat import b
from binascii import b2a_hex

fg = FortunaGenerator.AESGenerator()
fg.reseed(b(seedstring))
# Initial key is computed as
# SHAd256(b '\x00\x00...\x00'+b(seedstring))
print(fg.key)
fg.counter.next_value() # --> 1
b2a_hex(fg.pseudo_random_data(16))
print(fg.key)
fg.counter.next_value() # --> 4
```

# Generating random bytes in Linux

- Under Linux the user may extract high-quality cryptographic random bytes from the special files `/dev/random` and `/dev/urandom`
- Both files (actually, a character special files) harvest entropy from the same sources
- The fundamental difference is that requests to one of these (i.e., `/dev/random`) are *blocking*, while requests to the other are always immediately honored
- This means that if the generator has not sufficient entropy to satisfy the request, then the process that issued it is put on hold
- Under some circumstances, this might be preferable

# Using `/dev/random` and `/dev/urandom` from command shell

- The following command reads 24 bytes of (pseudo)random data from `/dev/urandom` and prints it to the terminal using `base64` encoding (for readability)

```
dd if=/dev/urandom bs=8 count=3 2>/dev/null|base64
```

(redirection of `stderr` is required to keep the output clean)

- When using `/dev/random` it may be useful to keep an eye on the available entropy beforehand

```
cat /proc/sys/kernel/random/entropy_avail
```

- ... and then using `/dev/random` under timeout guard

```
timeout 5s dd if=/dev/random/ bs=....
```

to avoid blocking (while possibly getting less data)

# PRNGs in Python

```
from os import urandom
from Crypto.Random.OSRNG import os
from binascii import b2a_hex

urandom is os.urandom      # eval to True
r = b2a_hex(urandom(32))   # returns 32 bytes
                           # of random data
```